

portals

The Production Memory Field Guide

the production repository for ai-native
creative organizations

The Production Memory Field Guide

How AI-native teams preserve the versions, context, decisions, and relationships behind their most valuable creative work

AI has made creative generation faster, cheaper, and more abundant. A production team can now create hundreds of images, shots, campaign variations, character studies, and visual directions in the time previously required to produce a much smaller body of work. But the cost of understanding that work has not kept pace. The final asset may be easy to store. The production state behind it is not: which version was approved, which prompt and settings produced it, which decisions shaped it, and how another contributor should reproduce or extend it. This field guide introduces production memory—the complete, recoverable organizational record of how an important asset was created—and provides a framework for preserving, evaluating, and recovering that record as creative production scales.

Contents

Executive Summary

The Problem with Production Memory

Risk 1 — Version Confusion

Risk 2 — Context Drift

Risk 3 — Decision Erosion

Risk 4 — Asset Orphanhood

Risk 5 — Handoff Failure

Risk 6 — Knowledge Attrition

The Cost of Broken Production Memory

Minimum Viable Production Memory Standard

Diagnostics — How to Assess Your Current State

Evaluation Framework — Manual vs. Software Solutions

When to Build a Production Repository

Implementation Guide — Getting Started with Production Memory

The Role of Portals — A Production Repository for AI-Native Teams

Appendix — Glossary of Terms

Appendix — Diagnostic Worksheet

Appendix — Recommended Reading

Appendix — About Portals

Executive Summary

AI has made creative generation faster, cheaper, and more abundant.

A production team can now create hundreds of images, shots, campaign variations, character studies, environments, voice treatments, and visual directions in the time previously required to produce a much smaller body of work.

But the cost of understanding that work has not fallen at the same rate.

The final asset may be easy to store. The production state behind it is not:

which version was approved;

which prompt, model, and settings produced it;

which references shaped it;

which decisions changed it;

which alternatives were rejected;

which assets it came from;

which derivatives came after it;

how another contributor should reproduce or extend it.

This creates a fundamental asymmetry in AI production:

“Generation is becoming inexpensive. Reconstruction remains expensive.”

The more a team generates, the more production context it creates. Unless that context is preserved deliberately, the organization accumulates assets faster than it accumulates understanding.

This field guide introduces the concept of production memory:

Production memory is the complete, recoverable organizational record of how an important asset was created: the version that was approved, the context that produced it, the decisions that shaped it, and the relationships between it and the work around it.

Production memory allows a team to:

identify the correct version;

understand how it was produced;
reproduce important conditions;
create controlled extensions;
preserve continuity;
transfer work between people;
trace variants back to their source;
retain production knowledge as an organizational asset.

This guide will help your team:

1. understand why production-memory failures are structural to AI-native work;
2. assess the current state of your production process;
3. recognize six recurring workflow risks;
4. quantify the cost of lost production context;
5. establish a minimum production-memory standard;
6. evaluate manual and software-based solutions;
7. determine when a dedicated production repository becomes necessary.

You do not need Portals to begin improving production memory.

Every team can make progress through clearer version designation, better context logging, stronger approval records, linked references, and disciplined handoff practices.

The objective is not perfect documentation.

The objective is recoverability:

“Can another qualified person find, understand, reproduce, and extend the work without depending on undocumented memory?”

In traditional creative production, memory is carried by people and artifacts.

The director remembers which take was best. The editor remembers which cut got approval. The art director remembers where the reference came from. The producer remembers why the team chose one VFX vendor over another. The file name carries a hint, the folder structure carries an implicit logic, and the review thread carries a compressed record of decisions.

These mechanisms are imperfect, but they have worked well enough because production was slow enough to leave traces. A live-action shoot produces a finite number of takes. A hand-drawn animation produces a finite number of frames. Each frame takes enough time that someone has time to think before the next one arrives.

AI production breaks all of these assumptions at once.

The speed of generation means that models, prompts, seeds, settings, and variants change faster than any human can track.

The volume of output means that what would have been a single asset in a traditional pipeline becomes hundreds or thousands of candidates.

The opacity of the model means that the cause of a result the conditions that led to a specific image, frame, or treatment is often invisible even to the person who generated it.

The iteration model means that the most important work often happens through small modifications to a latent space rather than through explicit decisions about a visible artifact.

These four factors speed, volume, opacity, and latent iteration create a structural memory problem.

The production record fragments the moment it is created. The artifacts persist. The context does not.

This is not a failure of process or discipline. It is a structural property of how AI production works. The traditional mechanisms for preserving production context were designed for a different pace and pattern of work. They cannot keep up.

The fragmentation of production memory creates specific, predictable problems that compound over time. The next six sections describe each of these risks in detail.

The Problem with Production Memory

In traditional creative production, memory is carried by people and artifacts.

The director remembers which take was best. The editor remembers which cut got approval. The art director remembers where the reference came from. The producer remembers why the team chose one VFX vendor over another. The file name carries a hint, the folder structure carries an implicit logic, and the review thread carries a compressed record of decisions.

These mechanisms are imperfect, but they have worked well enough because production was slow enough to leave traces. A live-action shoot produces a finite number of takes. A hand-drawn animation produces a finite number of frames. Each frame takes enough time that someone has time to think before the next one arrives.

AI production breaks all of these assumptions at once.

The speed of generation means that models, prompts, seeds, settings, and variants change faster than any human can track.

The volume of output means that what would have been a single asset in a traditional pipeline becomes hundreds or thousands of candidates.

The opacity of the model means that the cause of a result the conditions that led to a specific image, frame, or treatment is often invisible even to the person who generated it.

The iteration model means that the most important work often happens through small modifications to a latent space rather than through explicit decisions about a visible artifact.

These four factors speed, volume, opacity, and latent iteration create a structural memory problem.

The production record fragments the moment it is created. The artifacts persist. The context does not.

This is not a failure of process or discipline. It is a structural property of how AI production works. The traditional mechanisms for preserving production context were designed for a different pace and pattern of work. They cannot keep up.

The fragmentation of production memory creates specific, predictable problems that compound over time. The next six sections describe each of these risks in detail.

Risk 1 — Version Confusion

When a team generates 50 variants of a single shot in a single session, which one is the right one? The question sounds trivial. In practice, it is the source of more production friction than any other single failure.

Version confusion occurs when a team cannot reliably determine which asset is the approved, current, canonical version. It is the most common production-memory failure and the one with the most immediate consequences.

In a traditional workflow, versioning is relatively straightforward. A file name like `spot_v03_final_v2_approved` conveys meaningful information because the number of versions is small and the production history is comprehensible.

In AI production, the version tree explodes. A single prompt can generate 10 variations. Each variation can be iterated on with different settings. Each iteration can fork into multiple directions. Within hours, a team can have hundreds of assets with no reliable way to distinguish the approved version from experimental dead ends, intermediate drafts, or discarded alternatives.

The consequences of version confusion include: teams working from the wrong asset; rework that was completed but not propagated; decision reversals because the approved version cannot be found; time lost searching for the right file; and erosion of confidence in the production process itself.

Version confusion is not a discipline problem. It is a structural problem. The volume and speed of AI production overwhelm the naming conventions, folder structures, and mental models that teams have relied on for decades. When a single session produces more assets than a traditional project produced in its entirety, the old systems break.

The cost of version confusion is compounded by the fact that AI-generated assets are often difficult to distinguish from one another at a glance. Two images with meaningfully different generation parameters may look nearly identical until inspected closely. The difference matters — a model update, a seed change, or a prompt refinement can have a significant effect on downstream usability — but the visible similarity creates a false sense of equivalence.

Resolving version confusion requires a system that makes the canonical version immediately identifiable, preserves the relationship between versions, and records the decision process that led to the selection of one version over others.

Pilot Exercise: Pick one asset produced in the last week. Can three team members independently identify the approved version in under 30 seconds? If not, your team is experiencing version confusion.

Risk 2 — Context Drift

Context drift occurs when the conditions that produced an asset become unrecoverable. The team knows what was approved, but no longer knows how it was created.

In AI production, the how is essential. Unlike traditional tools where the process is embedded in the tool itself — a particular brush stroke, a specific camera angle, a defined composite — AI generation is mediated by opaque models and volatile parameters. The same prompt can produce meaningfully different results across model versions, seed values, guidance scales, and sampling methods.

When context is lost, the team cannot:

reproduce the asset if it needs to be regenerated at a higher resolution or different format;

understand why a particular result emerged from a given set of inputs;

diagnose quality regressions when model updates change output behavior;

train new team members on the production approach;

audit production decisions for compliance or client requirements;

build on prior work without rediscovering the approach from scratch.

Context drift is particularly insidious because it happens gradually. A team that carefully logs prompts and settings in week one may be too busy to log them in week four. By week eight, the generation process has become a black box. The assets exist. The context does not.

The cost of context drift is cumulative. Every lost context increases the time required to produce the next asset, because the team must rediscover what it already learned. The knowledge walks out the door whenever a team member leaves, because the knowledge was never captured outside their head.

Context drift is often invisible to leadership until a crisis — a client requests a small revision to an asset whose production context has been lost, and the team finds itself regenerating from scratch, hoping to approximate the original result.

Pilot Exercise: Find an AI-generated asset from one month ago. Can you reproduce the exact generation parameters that produced it? If the answer requires memory rather than a recorded source of truth, your team is experiencing context drift.

Risk 3 — Decision Erosion

Decision erosion is the gradual loss of the rationale behind production choices. The team may know what was decided, but no longer knows why.

In traditional production, decisions are often embedded in the artifacts themselves. The edit tells the story. The layout embodies the design rationale. The composition reflects the directorial intent. The decision is visible in the work.

In AI production, the most important decisions happen before the artifact exists: which model to use, which prompt strategy to pursue, which seed range to explore, which aesthetic direction to prioritize. These decisions are invisible in the final asset. The image itself does not reveal why it looks the way it does.

Decision erosion creates two distinct problems:

First, it prevents learning. When the rationale for a decision is lost, the team cannot evaluate whether the decision was correct, apply the same reasoning to future work, or avoid repeating mistakes. Each project starts fresh, with no accumulated strategic knowledge.

Second, it produces decision reversals. When the same question arises weeks or months later, a different person with different context makes a different decision — not because the first decision was wrong, but because the reasoning behind it is gone. The result is inconsistent output, wasted work, and organizational whiplash.

Decision erosion is particularly dangerous for AI-native production because the decision space is so large. A team might make hundreds of substantive decisions in a single production day. Without a system for recording the rationale behind those decisions, the vast majority of production intelligence is lost within days.

The most valuable thing a team can preserve is not the asset. It is the decision process that produced the asset.

Pilot Exercise: Review the last major creative decision your team made. Is the rationale for that decision written down somewhere accessible to the whole team? Could a new team member arriving tomorrow understand why that decision was made?

Risk 4 — Asset Orphanhood

Asset orphanhood occurs when an asset exists but its relationships to other assets, to the production context, and to the project itself are lost. The asset is present but its place in the production ecosystem is unknown.

In AI production, no asset exists in isolation. Every image is derived from a prompt, which may be derived from a reference, which may be derived from a style guide, which may be derived from a creative brief. Every asset is connected to other assets — earlier iterations, parallel variants, downstream composites, derivative works.

Asset orphanhood breaks these connections. An orphaned asset cannot answer basic questions: Where did this come from? What was it used for? What came after it? What depended on it?

The consequences of orphanhood include:

duplicate work, because the team does not know what already exists;

missed reuse, because the team does not know what could be repurposed;

broken chains, because downstream assets cannot be updated when a source changes;

lost provenance, because the origin of an important asset cannot be traced.

Asset orphanhood scales with production volume. A team producing 100 assets a week can maintain reasonable relationships manually. A team producing 1,000 assets a day cannot. The connections fray and snap, and the production graph fragments.

The cost of orphanhood extends beyond immediate inefficiency. When assets cannot be reliably connected to their production context, the value of the entire asset library degrades. A library of orphaned assets is a graveyard of disconnected artifacts. It holds potential value that can never be realized because the relationships that would unlock that value have been lost.

Pilot Exercise: Take one approved production asset. Can you trace its lineage — what reference or prompt produced it, what variants were considered, what assets were produced from it? If the chain breaks at any point, that asset is partially orphaned.

Risk 5 — Handoff Failure

Handoff failure occurs when work moves between people, teams, or stages without sufficient production context to maintain continuity.

In a traditional production pipeline, handoffs are shaped by the tools and roles involved. The storyboard artist hands off to the animator. The animator hands off to the compositor. Each handoff is accompanied by a specific set of conventions, expectations, and artifacts that carry meaning.

In AI production, the handoff is often undefined or invisible. The person who generated the initial concepts may be the same person who refines them, but the work may span multiple sessions, tools, and mental models that are never externalized. Alternatively, the person who generates concepts may be entirely separate from the person who curates and refines them, with no standard interface between the two roles.

Handoff failure manifests as:

dropped context: the receiving team does not know what decisions have already been made;

repeated work: the receiving team regenerates what was already produced because they cannot find or use it;

misaligned output: the receiving team produces work that does not match the intent because the intent was not communicated;

friction and delay: the handoff itself requires extensive meetings, documentation, and clarification that erode the speed advantage of AI production.

Handoff failure is a structural problem in AI production because the handoff often involves asymmetrical context. The person who generated the assets has deep, undocumented knowledge of what worked, what did not, and why certain paths were chosen. The person receiving those assets has only the assets themselves. The asymmetry is a risk.

The goal of a good handoff is not to transfer every detail of the production process. It is to transfer enough context that the receiving party can make competent decisions about the work without needing to reconstruct the sending party's production state.

Pilot Exercise: Identify the last handoff in your production process. Ask the person who received the work: Did you have everything you needed to understand the state of the work and proceed without clarification? If the answer is no, that handoff failed.

Risk 6 — Knowledge Attrition

Knowledge attrition is the loss of production capability that occurs when team members leave, roles change, or time passes without the organizational knowledge being preserved in a recoverable form.

In traditional production, knowledge attrition is mitigated by the fact that production knowledge is often embedded in institutional processes, documented workflows, and established toolchains. A new editor joining a post-production house can learn the Avid workflow from a manual. A new animator can study the rigging guide.

In AI production, much of the knowledge is tacit. It lives in the prompts that someone discovered through trial and error. In the settings that someone dialed in over weeks of experimentation. In the aesthetic judgments that someone developed through hundreds of iterations. This knowledge is rarely written down and almost never systematized.

The rate of knowledge attrition is accelerating because the field itself is changing so quickly. A production technique that was state of the art six months ago may be obsolete. A model that produced reliable results may have been superseded. The half-life of AI production knowledge is measured in months, not years.

Knowledge attrition has a compounding effect. When a team loses a member who held critical production knowledge, the remaining team not only loses that person's expertise — they also lose the ability to build on that knowledge. The production capability of the team degrades nonlinearly.

The most common response to knowledge attrition is to treat it as an HR problem: document processes, write guides, conduct exit interviews. These measures are valuable, but they miss the fundamental point. Knowledge in AI production is not primarily about process. It is about the specific, contextual, experiential knowledge of what works and why. That knowledge cannot be captured in a process document because it is too granular and too rapidly evolving.

What is needed is a system that preserves production knowledge as a natural byproduct of the production process itself rather than as a separate documentation activity that will be deprioritized under deadline pressure.

Pilot Exercise: Ask each team member to write down the five most important things they have learned about your AI production process in the last month. Compare the lists. If there is minimal overlap, your team is vulnerable to knowledge attrition.

The Cost of Broken Production Memory

The six risks described above are not theoretical. They have measurable operational costs that affect the speed, quality, and sustainability of AI production.

The most visible cost is rework. When a team cannot find the approved version, they regenerate it. When they cannot reproduce the conditions, they rediscover them. When they cannot trace the lineage, they retrace it. Every instance of lost production memory translates directly into time spent recreating what already existed.

A team that spends 10 percent of its production time on rework due to memory failures loses 25 productive days per person per year. For a team of 10, that is 250 days of lost capacity.

The less visible costs are often larger:

Missed reuse. When a team does not know what exists in its own library, it creates duplicates rather than repurposing existing work. Every duplicate is a wasted generation.

Reduced iteration. When the cost of understanding the current state is high, teams iterate less. They settle for good enough because exploring the decision space is too expensive.

Slower onboarding. New team members cannot ramp up by studying the production record because the production record does not exist. They must learn through trial and error, repeating the same discovery process that previous team members already completed.

Strategic blind spots. Without accumulated production knowledge, teams cannot identify patterns in their own work. They do not know which approaches are consistently effective, which models produce the most reliable results, which prompt strategies yield the highest approval rates.

The aggregate cost of broken production memory is often 2-3x the direct labor cost of rework. It is a tax on every aspect of production that compounds as the team grows, the project expands, and the production history accumulates.

Organizations that track these costs consistently find that the investment required to preserve production memory is a fraction of the cost of losing it.

Minimum Viable Production Memory Standard

Before evaluating any tool or system, a team should establish a minimum standard for what production memory means in practice. The standard does not need to be comprehensive. It needs to be consistent.

The Minimum Viable Production Memory Standard consists of five requirements:

1. Every approved asset must have an unambiguous, findable version identifier.

The version identifier is the atomic unit of production memory. Without a reliable way to designate which version of an asset is the canonical one, every other production-memory practice is built on sand.

The identifier can be a file name convention, a database entry, a tag in a DAM system, or a field in a production-tracking spreadsheet. The format matters less than the consistency. Every approved asset must have exactly one canonical version identifier, and it must be findable by everyone on the team.

2. Every approved asset must have a recorded production context.

The minimum production context includes: the model and version used, the prompt or input that generated the asset, the key settings (seed, guidance scale, sampling method, resolution), and the date of generation.

This is the bare minimum. Teams that can capture more context particularly the rationale for key decisions will build a stronger production memory over time.

3. Every approved asset must have a recorded approval record.

The approval record must include: who approved the asset, when it was approved, and what was approved (the specific version, not the general direction). If an approval was conditional, the conditions must be recorded.

4. Every approved asset must have recorded relationships to its source and derivative assets.

At minimum: what previous version or reference this asset was derived from, and what subsequent versions or composites were produced from this asset.

These relationships form the production graph. Without them, every asset is an orphan.

5. The standard must be enforced consistently, not aspirationally.

A standard that is applied 60 percent of the time is not a standard. It is an aspiration. The five requirements above are deliberately minimal so that a team can meet them consistently without significant overhead. A team that meets the minimum standard 100 percent of the time has a stronger production memory system than a team that meets an ambitious standard 60 percent of the time.

The minimum standard is achievable with manual processes, spreadsheets, and discipline. No software purchase is required. The purpose of the standard is to create a baseline of consistency before evaluating whether tooling is needed to maintain or scale that baseline.

Diagnostics — How to Assess Your Current State

Before deciding what to fix, a team needs to know what is broken. The following diagnostic framework provides a structured way to assess the current state of production memory across the six risk areas.

How to use this framework

For each risk area, rate your team on a scale of 1 to 5:

- 1 — We experience this risk constantly and it is causing significant problems.
- 2 — We experience this risk regularly and it is causing noticeable friction.
- 3 — We experience this risk occasionally but it is usually manageable.
- 4 — We rarely experience this risk and it is generally under control.
- 5 — We do not experience this risk. Our production memory is healthy here.

Diagnostic Questions

Version Confusion: Can three team members independently identify the approved version of any recent asset in under 30 seconds?

Context Drift: Can you reproduce the generation parameters for any asset produced more than two weeks ago?

Decision Erosion: Is the rationale for major creative decisions documented and accessible?

Asset Orphanhood: Can you trace the lineage of any approved asset back to its source and forward to its derivatives?

Handoff Failure: Does work move between people or stages with sufficient context to maintain continuity?

Knowledge Attrition: If a key team member left today, would their production knowledge remain accessible?

Scoring

30 — Perfect production memory (rare in practice).

20-29 — Strong production memory. Targeted improvements will close remaining gaps.

12-19 — Moderate production memory. Several risk areas need attention.

6-11 — Weak production memory. Systemic failures across multiple areas.

Below 6 — Critical production memory failures. The team is operating without a functional memory system.

After completing the diagnostic, identify the three lowest-scoring areas as your initial focus. Improvements in these areas will produce the highest return on investment.

Evaluation Framework — Manual vs. Software Solutions

Once a team has assessed its current state and established a minimum standard, the next question is whether to rely on manual processes, adopt software tooling, or use a combination of both.

The answer depends on several factors. The following framework helps teams evaluate their specific situation and make an informed decision.

When manual processes are sufficient

Manual processes can work well when: the team is small enough that informal communication covers most hand-offs; production volume is low enough that manual tracking is feasible; the team has strong discipline and low turnover; the production workflow is simple and linear; and the cost of memory failure is low enough that occasional rework is acceptable.

For teams in this situation, the minimum standard described above typically provides a sufficient framework. A shared spreadsheet, consistent file naming conventions, and regular production reviews can maintain adequate production memory without dedicated software.

When software tooling adds value

Software tooling becomes valuable when: the team grows beyond 5-8 people; production volume exceeds 50-100 assets per week; the production workflow branches into parallel streams; the team experiences regular handoffs between roles or stages; turnover creates knowledge gaps; the cost of rework is significant; or the team needs to audit production decisions for clients or compliance.

Software tooling is not a substitute for the minimum standard. It is an enabler of the standard at scale. A team that cannot meet the minimum standard manually will not meet it with software either.

Evaluation criteria

When evaluating production memory solutions, consider these criteria:

Capture friction — How much additional effort does the solution require from producers? The best solutions capture context automatically as a byproduct of the production workflow.

Retrieval speed — How quickly can a team member find the production context for any asset? Seconds, not minutes.

Relationship preservation — Does the solution capture connections between assets, or does it treat each asset in isolation?

Team accessibility — Can everyone on the team access and contribute to production memory, or is it gated by tool access or expertise?

Integration — Does the solution work with the tools the team already uses, or does it require switching to a new workflow?

Scalability — Does the solution continue to work as production volume and team size grow?

When to Build a Production Repository

For many teams, the minimum standard and manual processes are sufficient. For others, the volume and complexity of AI production demand a dedicated production repository a system designed specifically to capture, preserve, and make accessible the complete production record.

A production repository is distinct from a DAM or asset management system. A DAM stores finished assets. A production repository stores the production context behind those assets: the versions, the prompts, the settings, the decisions, the relationships, the approvals, and the history.

The threshold for needing a production repository is crossed when:

- the team cannot maintain the minimum standard through manual processes alone;

- the cost of rework from memory failures exceeds the cost of the repository;

- the production graph has become too complex to track manually;

- the team needs to audit production decisions at scale;

- the organization treats production knowledge as a strategic asset worth preserving.

A production repository should:

- capture production context automatically from the tools the team already uses;

- make that context searchable, traversable, and recoverable;

- preserve the relationships between assets, decisions, and approvals;

- provide a single source of truth for production state;

- integrate with existing tools and workflows rather than replacing them.

The right time to build or buy a production repository is before memory failures become the bottleneck in your production process, not after. Waiting until the problem is acute means the team has already lost production knowledge that will never be recovered.

Implementation Guide — Getting Started with Production Memory

Improving production memory does not require a major initiative. It requires starting. The following implementation guide provides a phased approach that any team can adapt to its specific situation.

Phase 1: Diagnose and set a baseline (Week 1)

Complete the diagnostic assessment. Establish a baseline score across the six risk areas. Share the results with the team to build awareness of production memory as a discipline.

Output: A baseline diagnostic score and a shared understanding of the teams production memory strengths and weaknesses.

Phase 2: Implement the minimum standard (Weeks 2-3)

Adopt the five requirements of the Minimum Viable Production Memory Standard. Choose the simplest possible implementation for each requirement. Use existing tools wherever possible. The goal is consistency, not sophistication.

Output: A functioning minimum standard applied to all new production work.

Phase 3: Address the top three risks (Weeks 4-6)

Based on the diagnostic results, identify the three most impactful risk areas. For each area, implement one targeted improvement. An improvement does not need to be comprehensive. It needs to move the needle.

Output: Measurable improvement in the three prioritized risk areas.

Phase 4: Evaluate tooling needs (Week 7)

After six weeks of consistent application, reassess the diagnostic. Evaluate whether manual processes are sufficient or whether software tooling would add value. Use the evaluation framework to make an informed decision.

Output: A clear decision on whether to adopt production-memory tooling.

Phase 5: Scale and sustain (Ongoing)

Production memory is not a project with an end date. It is an ongoing operational discipline. Re-assess the diagnostic quarterly. Review the minimum standard for relevance. Adapt the approach as the team, tools, and production volume evolve.

Output: A sustainable production memory practice that evolves with the team.

The Role of Portals — A Production Repository for AI-Native Teams

Portals is a production repository designed specifically for AI-native creative organizations. It is built on the premise that production memory is the most undervalued asset in AI production.

Portals addresses each of the six risks described in this guide:

Version Confusion: Portals makes every version identifiable, comparable, and retrievable. The approved version is always marked and always findable.

Context Drift: Portals automatically captures generation parameters, model information, and production context from integrated tools. Context is preserved by default, not by effort.

Decision Erosion: Portals records decisions alongside assets, preserving the rationale as well as the result. Decision history is traversable and auditable.

Asset Orphanhood: Portals maintains the production graph automatically, tracking relationships between assets, versions, prompts, references, and derivatives.

Handoff Failure: Portals provides a shared production state that every team member can access, reducing the context asymmetry that causes handoff failures.

Knowledge Attrition: Portals preserves production knowledge as a natural byproduct of the production process, making it independent of any individual team member.

Portals is not a replacement for creative tools. It is a layer that captures and preserves the production context those tools generate, making it accessible, searchable, and actionable across the organization.

Portals works with the tools your team already uses: Midjourney, DALL-E, Stable Diffusion, Runway, ElevenLabs, and the growing ecosystem of AI creative tools. It integrates through APIs, plugins, and manual capture to create a unified production record.

If your team has completed the diagnostic, implemented the minimum standard, and found that manual processes cannot keep pace with production volume, Portals provides a purpose-built solution for preserving production memory at scale.

Appendix — Glossary of Terms

Approved version — The canonical version of an asset that has been reviewed and accepted for use. The single source of truth for production state.

Asset lineage — The chain of relationships tracing an asset back to its source materials and forward to its derivatives.

Context drift — The gradual loss of the production conditions that produced an asset.

Decision erosion — The loss of the rationale behind production choices over time.

Generation parameters — The specific settings used to produce an AI-generated asset, including model, prompt, seed, guidance scale, sampling method, and resolution.

Knowledge attrition — The loss of production capability that occurs when team members leave or organizational knowledge degrades.

Production graph — The network of relationships connecting assets, versions, decisions, approvals, and production context across a project or organization.

Production memory — The complete, recoverable organizational record of how an important asset was created: the version that was approved, the context that produced it, the decisions that shaped it, and the relationships between it and the work around it.

Production repository — A system designed specifically to capture, preserve, and make accessible the complete production record of an organization.

Version confusion — The inability to reliably determine which asset is the approved, current, canonical version.

Appendix — Diagnostic Worksheet

This worksheet provides a structured format for conducting the production memory diagnostic with your team.

Instructions

For each risk area, rate your team from 1 to 5. Record evidence for each score. Identify the three lowest-scoring areas as improvement priorities.

Version Confusion

Score (1-5): ___

Evidence: ___

Context Drift

Score (1-5): ___

Evidence: ___

Decision Erosion

Score (1-5): ___

Evidence: ___

Asset Orphanhood

Score (1-5): ___

Evidence: ___

Handoff Failure

Score (1-5): ___

Evidence: ___

Knowledge Attrition

Score (1-5): ___

Evidence: ___

Total Score

Sum of all six scores: ___ / 30

Priority Areas

1. (lowest score): ___

2. (second lowest): ___

3. (third lowest): ___

Action Items

For each priority area, define one specific improvement to implement in the next two weeks.

Appendix — Recommended Reading

These resources provide deeper context on the concepts introduced in this field guide.

Knowledge Management

The Knowledge-Creating Company — Ikujiro Nonaka and Hirotaka Takeuchi. The foundational text on how organizations create and preserve knowledge.

Working Knowledge — Thomas H. Davenport and Laurence Prusak. A practical guide to how organizations manage what they know.

Version Control and Lineage

The Git Book — Scott Chacon and Ben Straub. While designed for code, Git's approach to versioning, branching, and merging offers valuable lessons for creative production.

AI Production and Creative Operations

The AI Production Handbook — Various contributors. Practical guidance on building AI production workflows.

Portals Blog (portals.works/blog) — Case studies and best practices for production memory in AI-native creative organizations.

Appendix — About Portals

Portals is a production repository for AI-native creative organizations. We help teams preserve, navigate, and recover the production context behind their most valuable creative work.

Our platform captures generation parameters, tracks versions, records decisions, maintains the production graph, and makes production memory accessible across the organization.

Portals integrates with the tools AI-native teams already use: Midjourney, DALL-E, Stable Diffusion, Runway, ElevenLabs, and the growing ecosystem of AI creative tools.

Portals is built by a team of producers, engineers, and designers who have experienced firsthand the cost of broken production memory in AI-native workflows. We built Portals because we needed it ourselves.

To learn more or scope a production pilot, visit portals.works or contact sales@portals.works.